

# High Performance Computing: Homework 2

Dmitriy (Tim) Kunisky [dk3105]

- Processor: Intel(R) Core(TM) i5-5257U CPU @ 2.70GHz
  - Maximum flop rate: 19.76 Gflop/s, or 4.94 Gflop/s per core
  - Maximum main memory bandwidth: 25.6 GB/s
- Compiler: Apple LLVM version 9.1.0 (clang-902.0.39.2)

## Problem 2: Matrix-Matrix Multiplication

**Loop orderings.** The fastest loop order appears to be  $(j, p, i)$ , corresponding to traversing columns with temporal locality as much as possible (since traversing the  $i$  index traverses a column of both  $\mathbf{A}$  and  $\mathbf{C}$ , traversing the  $p$  index traverses a column of  $\mathbf{B}$  and a row of  $\mathbf{A}$ , and traversing the  $j$  index traverses a row of both  $\mathbf{B}$  and  $\mathbf{C}$ ). We then expect the most harmful changes to this ordering to be those that move column traversals earlier in the loop order. Thus, the costliest ordering should be the reverse,  $(i, p, j)$ , while orderings like  $(j, i, p)$  should be intermediate. The table below illustrates that this is indeed the case.

| Dimension | Ordering $(j, p, i)$   |                     | Ordering $(j, i, p)$   |                     | Ordering $(i, p, j)$   |                     |
|-----------|------------------------|---------------------|------------------------|---------------------|------------------------|---------------------|
|           | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) |
| 16        | 5.93                   | 94.82               | 2.53                   | 40.51               | 2.62                   | 41.94               |
| 112       | 9.62                   | 154.00              | 1.83                   | 29.23               | 1.54                   | 24.70               |
| 304       | 7.79                   | 124.64              | 0.91                   | 26.98               | 0.77                   | 12.34               |
| 640       | 5.66                   | 90.64               | 0.82                   | 13.15               | 0.65                   | 10.43               |
| 976       | 4.38                   | 70.05               | 0.87                   | 13.90               | 0.38                   | 6.05                |
| 1408      | 4.17                   | 66.71               | 0.69                   | 11.12               | 0.19                   | 3.07                |

**Block size tuning.** It appears that the optimal block size is around 60. One possible justification for this is that the processor used for these tests has a Level 1 cache of size 64KB. If the block size is  $B$ , then the number of bytes required to store two  $B \times B$  arrays of doubles (the size of the data on which arithmetic is being performed in one block matrix operation) is  $2 \cdot 8 \cdot B^2$ , which is smaller than 64KB only while  $B \leq 60$  (over values of  $B$  divisible by 4). The figures in the table below support this heuristic calculation.

| Dimension | $B = 40$               |                     | $B = 60$               |                     | $B = 80$               |                     |
|-----------|------------------------|---------------------|------------------------|---------------------|------------------------|---------------------|
|           | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) |
| 240       | 4.69                   | 75.05               | 8.95                   | 143.26              | 9.12                   | 145.91              |
| 480       | 5.37                   | 85.93               | 8.79                   | 140.78              | 7.92                   | 126.78              |
| 720       | 5.28                   | 84.55               | 8.68                   | 138.85              | 8.38                   | 134.07              |
| 960       | 5.26                   | 84.20               | 8.80                   | 140.82              | 6.52                   | 104.37              |
| 1200      | 5.41                   | 86.48               | 8.33                   | 133.21              | 6.84                   | 109.41              |
| 1440      | 5.19                   | 82.97               | 8.56                   | 137.01              | 7.44                   | 119.04              |
| 1680      | 5.31                   | 85.03               | 8.16                   | 130.55              | 6.87                   | 109.91              |

**Parallelization.** Finally, we parallelize our implementation by simply using the parallelized for loop directive of OpenMP on the top-level block iteration. This roughly doubles the flop rate, which corresponds to this processor having two cores. With this optimization, the matrix multiplication reaches a flop rate of slightly more than 17 Gflop/s, which is approximately 86% of the peak flop rate of 19.76 Gflop/s.

| Dimension | $B = 60$ , Serial      |                     | $B = 60$ , Parallel    |                     |
|-----------|------------------------|---------------------|------------------------|---------------------|
|           | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) | Flop Rate<br>(Gflop/s) | Bandwidth<br>(GB/s) |
| 240       | 8.95                   | 143.26              | 19.68                  | 314.86              |
| 480       | 8.79                   | 140.78              | 18.41                  | 310.60              |
| 720       | 8.68                   | 138.85              | 18.68                  | 298.84              |
| 960       | 8.80                   | 140.82              | 17.01                  | 272.16              |
| 1200      | 8.33                   | 133.21              | 17.85                  | 285.60              |
| 1440      | 8.56                   | 137.01              | 17.59                  | 281.39              |
| 1680      | 8.16                   | 130.55              | 17.24                  | 275.79              |

#### Problem 4: Jacobi/Gauss-Seidel Smoothing

We present timing results for varying problem sizes of the Jacobi and Gauss-Seidel smoothers, for one, two, and four threads. The machine used has two cores, so we expect to see a benefit from two threads versus one, but no additional benefit from four threads versus two. The results below support this claim; for large problems, the speed roughly doubles from moving from one to two threads, but there is no effect from moving from two to four threads.

| N   | Jacobi   |           |           | Gauss-Seidel |           |           |
|-----|----------|-----------|-----------|--------------|-----------|-----------|
|     | 1 Thread | 2 Threads | 4 Threads | 1 Thread     | 2 Threads | 4 Threads |
| 50  | 0.15     | 0.19      | 0.18      | 0.25         | 0.20      | 0.28      |
| 100 | 0.55     | 0.34      | 0.35      | 0.90         | 0.57      | 0.65      |
| 150 | 1.35     | 0.68      | 0.84      | 2.03         | 1.10      | 1.22      |
| 200 | 2.28     | 1.13      | 1.51      | 3.54         | 2.01      | 2.20      |
| 300 | 5.03     | 2.70      | 3.39      | 7.89         | 4.18      | 4.38      |
| 400 | 10.46    | 7.86      | 7.64      | 17.31        | 7.74      | 8.05      |

(Note that running the parallel version of either algorithm is implemented by passing an arbitrary third argument to the executable, as shown in the top comment of either source file.)